

Akash Kanaujiya

AI Engineer | Backend Engineer | Distributed Systems Developer

Targeting to leverage expertise in AI-powered Back-end engineering, distributed systems, and scalable cloud architectures to drive the development of next-generation intelligent products with measurable business impact.

+91-7007694698

iakanaujiya@gmail.com

Gurugram, India

[LinkedIn](#)

[Leetcode](#)

[Github](#)



PROFILE SUMMARY

- Software Engineer with nearly 2 years of experience in building scalable Back-end platforms, AI applications, and distributed cloud-native systems using Python, Java, AWS, and modern data architectures.
- Currently working as Software Engineer 2 at HashedIn by Deloitte, delivering enterprise-grade AI solutions, real-time communication platforms, autonomous agents, and high-performance microservices for large-scale business use cases.
- Recognized with the Top Impactor Award for architecting and delivering a scalable AI ecosystem that improved operational efficiency, accelerated hiring workflow, and transformed enterprise support processes.
- Proven expertise in Agentic AI, RAG architectures, LLM orchestration, event-driven systems, and cloud-native application development, enabling organizations to scale AI adoption while optimizing business productivity and user engagement.
- Combines strong software engineering fundamentals with product-centric thinking, translating complex business challenges into scalable, secure, and outcome-driven technology solutions.



WORK EXPERIENCE

Software Engineer 2

HashedIn by Deloitte | Oct'24-Present

Gurugram, India

- Driving end-to-end design and development of AI-powered interview intelligence platforms supporting real-time hiring workflow.
- Orchestrating autonomous AI screening solutions to streamline candidate evaluation and recruitment operations.
- Designing scalable event-driven microservices leveraging Redis Streams, Pub/Sub, AWS services, and WebSocket infrastructure.
- Developing enterprise-grade RAG applications and conversational AI solutions for internal business operations.
- Collaborating across multi-technology ecosystems spanning Java, Python, C#, and TypeScript to accelerate product delivery.
- Building large-scale code intelligence and developer productivity platforms for enterprise engineering teams.

Achievements:

- Architected a real-time AI Interview Co-pilot platform that reduced interviewer effort by 75%, improved hiring decision accuracy by 50%, and supported 50+ concurrent live interview sessions through AWS-powered microservices.
- Implemented an AI-driven Autonomous Screening Agent that decreased manual candidate evaluation efforts by 80% while notably accelerating recruitment turnaround times.
- Established a horizontally scalable event-driven ecosystem that improved real-time communication responsiveness by 40% and enabled 3X expansion in WebSocket connectivity without performance degradation.
- Revolutionized employee support operations through a RAG-powered HR assistant, improving query resolution accuracy by 48%, reducing support tickets by 90%, and generating 20+ hours of weekly operational savings.
- Developed an automated code intelligence platform capable of analyzing 1M+ lines of multi-language code, reducing documentation effort by 90% and accelerating developer onboarding by 80%.



PROJECTS

Realtime AI Voice Chat

Python Asyncio, FastAPI, WebSockets, LangChain, Silero VAD

[Source Code](#) [Demo Video](#)

- Engineered a real-time voice communication platform enabling seamless integration with existing chatbot agents.
- Designed a modular orchestration framework supporting multiple Speech-to-Text (STT), Text-to-Speech (TTS), and Large Language Model (LLM) providers, reducing integration effort for new providers by approximately 80%.
- Optimized asynchronous audio streaming pipelines using Python Asyncio and WebSockets to achieve sub-second response latency and improve conversational responsiveness.
- Developed full-duplex audio processing workflow for simultaneous speech input and output, enabling natural real-time conversations.



CORE COMPETENCIES

- Agentic AI Development
- LLM Orchestration Frameworks
- Retrieval-Augmented Generation (RAG)
- AI Evaluation Frameworks
- Distributed Systems Engineering
- Event-Driven Architecture
- Cloud-Native Application Design
- Enterprise Microservices Architecture
- Real-Time Communication Systems
- Vector & Graph Database Solutions
- Context Engineering & Memory Systems
- AI Product Engineering



TECHNICAL SKILLS

- Programming Languages:** Python, Java, C#, C++, JavaScript, TypeScript, SQL, HTML5, CSS3
- Backend Technologies:** FastAPI, Django, Spring Boot, ASP.NET/.NET, REST APIs, WebSockets, AsyncIO
- Frontend Technologies:** React.js, Next.js, Tailwind CSS
- AI & GenAI Technologies:** LangChain, LangGraph, OpenAI APIs, Anthropic APIs, RAG, Vector Databases, Graph Databases, Agent Memory, Prompt Engineering, Context Engineering, STT, TTS
- Databases:** PostgreSQL, Aurora PostgreSQL, MongoDB, Redis
- Cloud & Infrastructure:** AWS Lambda, AWS ECS, AWS EC2, AWS S3, AWS SQS, AWS VPC, Docker
- Messaging & Processing:** Redis Pub/Sub, Redis Streams, Celery, SQLAlchemy
- Testing & Development Tools:** PyTest, JMeter, Git, GitHub, Linux, Visual Studio Code



ACHIEVEMENTS

- Awarded Top Impactor Award at HashedIn by Deloitte for delivering a high-impact AI solution.
- Ranked in the Top 4% globally on LeetCode with a 1900+ Knight rating. [Profile](#)
- Secured Global Rank 350 among 42K+ participants in a LeetCode Biweekly Contest.
- Attained Global Rank 513 in TCS CodeVita among 444K+ participants across 94 countries.
- Earned Meta Full-Stack Development with Django Certification (Coursera).
- Obtained JavaScript Security Specialization Certification.
- Claude Certified Architect

- Implemented event-driven processing pipelines to improve scalability, maintainability, and extensibility of AI service integrations.
- Improved overall user experience through latency optimization and efficient audio processing strategies

CodeRunr

FastAPI, Docker, PostgreSQL, Celery, SQLAlchemy, AWS

[Source Code](#)

- Architected a distributed remote code execution platform capable of scaling to millions of submissions per day with a target system availability of 99.9%.
- Designed secure sandboxing and process isolation mechanisms using Linux namespaces and cgroups to safely execute untrusted user code.
- Developed asynchronous execution pipelines using FastAPI, Celery, and distributed AWS services to improve job throughput and execution efficiency.
- Integrated AWS Lambda, EC2, SQS, and Aurora PostgreSQL to build a scalable, cloud-native execution and scheduling infrastructure.
- Optimized execution workflow to achieve average code execution times of approximately five seconds while maintaining platform reliability.
- Built modular Back-end services and database layers using SQLAlchemy and PostgreSQL to support extensibility and maintainable system design.
- Improved system scalability, fault tolerance, and operational efficiency through distributed processing and cloud-native architecture



- B.Tech., Computer Science & Engineering | Ajay Kumar Garg Engineering College, Ghaziabad | Jun'24 | CGPA: 7.9/10
Relevant Coursework: DSA, Operating Systems, DBMS, Machine Learning, System Design.