

Akash Kanaujiya

iakanaujiya@gmail.com | +91-7007694698 | Gurugram, India
about.akanaujiya.com | [LeetCode](#) | [GitHub](#) | [LinkedIn](#)

Education

Ajay Kumar Garg Engineering College, Ghaziabad, U.P.
Bachelor of Technology (B.Tech.) in Computer Science & Engineering

Jun 2024
CGPA: 7.9/10

Technical Skills

Languages: Python, Java, C++, C#, JavaScript, TypeScript, SQL, HTML, CSS

Frameworks: FastAPI, Spring Boot, Django, Express.js

Developer Tools: AWS (Lambda, ECS, SQS, VPC, S3), Docker, Redis, PostgreSQL, MongoDB, Git, Linux

AI/ML: LangChain, LangGraph, RAG, Prompt Engineering, Context Engineering

Core: Data Structures & Algorithms, Operating Systems, Computer Networks, DBMS, OOPs, System Design

Experience

HashedIn by Deloitte | *Software Engineer 2*

Oct 2024 – Present | *Gurugram*

- Designed and developed scalable backend services for a real-time AI interview co-pilot using **Python, FastAPI, WebSockets, PostgreSQL, Redis, Langchain, LangGraph and AWS**, supporting more than **50 concurrent interview sessions**.
- Developed event-driven microservices using **Redis Streams, SQS, Pub/Sub, and Python AsyncIO**, improving processing throughput and supporting approximately **3× more concurrent WebSocket connections**.
- Built an **autonomous candidate-screening workflow** using **LangGraph agents** to conduct structured interviews, track question coverage, and generate candidate evaluations, reducing manual screening effort by approximately **80%**.
- Built a **Retrieval-Augmented Generation chatbot** using **LangChain, PostgreSQL/pgvector, and Elasticsearch**, enabling employees to retrieve information from internal knowledge sources through natural-language queries.
- Designed **context-engineering** pipelines for agentic workflows using dynamic retrieval, conversation summarization, context prioritization, and token-budget management to improve response quality and reduce unnecessary LLM usage.
- Integrated **Langfuse** to monitor LLM prompts, responses, latency, and token usage, enabling prompt optimization and cost reduction across agentic workflows.
- Developed a code-intelligence platform capable of processing large multilingual repositories to generate documentation, identify dependencies, and accelerate developer onboarding.
- Designed and maintained REST APIs and backend services using **Java, Spring Boot, PostgreSQL, and Docker**, integrating new features with existing enterprise systems.

Achievements

- LeetCode Knight:** Achieved a peak contest rating of **1912**, placing in the global top **4%**.
- Achieved **Global Rank 350 out of 42,000+ participants** in LeetCode Biweekly Contest 492.
- Secured **513th rank globally** in TCS CodeVita, competing against more than **444,000 participants** across 94 countries.
- Received the **Top Impactor Award** at HashedIn by Deloitte for architecting and delivering a scalable AI system.
- Claude Certified Architect:** Obtained certification in Agentic AI from Anthropic.

Project Work

Real-Time AI Voice Chat | *Python AsyncIO, FastAPI, WebSockets, LangChain, Silero VAD*

[Source Code](#) | [Demo Video](#)

- Engineered a **real-time AI voice conversation system** that integrates with existing chatbot agents and transforms them into interactive voice-based interfaces.
- Designed a modular orchestration framework supporting multiple Speech-to-Text, Text-to-Speech, and Large Language Model providers, reducing integration effort for new providers by approximately **80%**.
- Optimized asynchronous full-duplex audio pipelines using Python AsyncIO and WebSockets to achieve sub-second response latency and improve conversational responsiveness.

CodeRunr (Remote Code Execution Engine) | *FastAPI, Docker, PostgreSQL, Celery, SQLAlchemy, AWS*

[Source Code](#)

- Built a remote code-execution platform using FastAPI, Celery, AWS SQS, and EC2, with isolated execution environments for untrusted programs.
- Implemented asynchronous job processing, execution timeouts, resource limits, and retry handling to improve reliability under concurrent workloads.
- Deployed the platform using AWS Lambda, SQS, EC2, and PostgreSQL, separating API and execution workloads for independent scaling.